

11-442 / 11-642 / 11-742: Search Engines


Index Creation

Jamie Callan
Carnegie Mellon University
callan@cs.cmu.edu

1

Lecture Outline

- Building inverted lists on a single processor
- Inverted lists and inverted files
 - Inverted list compression
 - Inverted list optimizations
- Forward indexes
- Index updates



2

© 2021, Jamie Callan

2

Basic Facts That Affect Indexing


Usually the corpus is much bigger than the available RAM

- E.g., 20 million web documents is 500 GB
- You can't do the whole task in memory

Disks are slow compared to processors

- Only use the disk when absolutely necessary
- Compress data to reduce I/O
- Sequential access is much faster than random access

Most of this is true for all parts of the search engine (e.g., 500 GB)
...but it is especially important during indexing



3

© 2021, Jamie Callan

3

Overview of Index Construction

The parser uses an API to communicate with the indexer

- There is no standard API
- We will take a quick look at what Lucene does

4

© 2021, Jamie Callan

4

Lucene's Indexing API: Configure the Inverted Index

Configure the lexical analyzer

```
EnglishAnalyzer analyzer = new EnglishAnalyzer ();
analyzer.setLowercase (true);
analyzer.setStopwordRemoval (true);
analyzer.setStemmer (EnglishAnalyzer.StemmerType.KSTEM);
```

Configure the inverted index

```
IndexWriterConfig iwc = new IndexWriterConfig(analyzer);
iwc.setOpenMode(OpenMode.CREATE);
IndexWriter writer = new IndexWriter(path, iwc);
```

5

© 2021, Jamie Callan

5


Lucene's Indexing API: Define Field Types

Define field types to control how different parts of the document are indexed

- Example: A field type for full-text fields (e.g., body, title)
 - Tokenize the content and store it in the inverted index and the forward index
 - Inverted lists contains docids, tf, positions
 - TermVectors (the forward index) contain positions

```
FieldType fullText = new FieldType();
fullText.setTokenized(true);
fullText.setIndexOptions(
    IndexOptions.DOCS_AND_FREQS_AND_POSITIONS);
fullText.setStoreTermVectors(true);
fullText.setStoreTermVectorPositions(true);
```

A document



© 2021, Jamie Callan

6

Lucene's Indexing API: Document Objects

Start a loop over all available documents

Create document object

```
Document doc = new Document();
```

7

© 2021, Jamie Callan

7

Lucene's Indexing API: Adding Fields to a Document Object

Create metadata field objects

```
f1 = new StringField("externalId", "GX016-79-782", Field.Store.YES);
```

```
f2 = new StringField("PageRank", "4.33", Field.Store.YES);
```

Create content field objects

```
f3 = new Field("title", "Juice Lyrics", fullText);
```

```
f4 = new Field("body", "It ain't my fault that I'm out here gettin' loose...", fullText);
```

Add field objects to the document object

- doc.add(f1);
- doc.add(f2);
- ...

8

© 2021, Jamie Callan

8

Lucene's Indexing API: Indexing a Document Object

Add the document to the index

- `IndexWriter.addDocument (doc);`


End the loop over all available documents

9

© 2021, Jamie Callan

9

Indexing Text Tokens



Most text representation tasks are done in the indexer

- Documents from different sources can be treated the same
- The indexer knows what data structures it needs

Typical text representation tasks (earlier lecture)

- Case folding (mixed case → lower case)
- Stopword removal
- Stemming (morphological processing / lemmatization)
- ...

10

© 2021, Jamie Callan

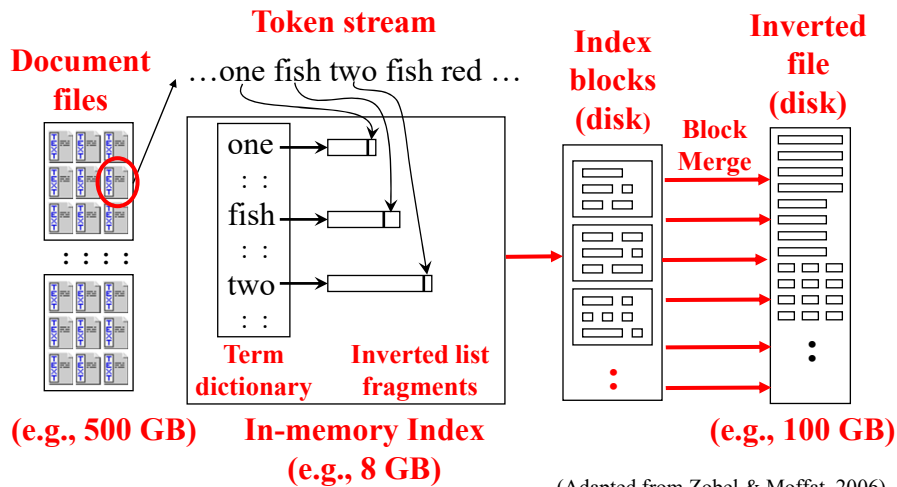
10

How Inverted Files are Built: Single Processor

Example entry for "stock"

Doc	id	value	date	price	volume
Doc1	1	8	8	1	1
Doc2	1	1	8	2	1
Doc3	1	1	1	1	1
Doc4	1	8	1	8	1

Inverted Lists



11

© 2021, Jamie Callan

11

How Inverted Files are Built: Single Processor

Example entry for "stock"

Doc	id	value	date	price	volume
Doc1	1	8	8	1	1
Doc2	1	1	8	2	1
Doc3	1	1	1	1	1
Doc4	1	8	1	8	1

Inverted Lists

The in-memory index buffers store

- Part of the term dictionary
- Fragments of inverted lists

The in-memory buffers are small compared to the final index

When in-memory buffers are full

- Flush in-memory buffers to disk and reinitialize them
- Continue parsing to refill the in-memory buffers

When all documents are parsed, merge index blocks on disk

- Very fast – essentially a merge of sorted lists

12

© 2021, Jamie Callan

12

Example entry for "stock"

Corpus

	a	abbe	about	ability	able	about	...	mean
Doc ₁	1	0	0	1	1	1
Doc ₂	1	1	0	0	1	1	...	0
Doc ₃	1	0	1	1	0	1	...	0

Inverted Lists



© 2021, Jamie Callan

14

Block merge of the inverted list for 'apple'



© 2021, Jamie Callan

Lecture Outline

- Building inverted lists on a single processor
- **Inverted lists and inverted files**
 - Inverted list compression
 - Inverted list optimizations
- Forward indexes
- Index updates

15

© 2021, Jamie Callan

15

Inverted List Indexes

Corpus Vocabulary

	a	alks	planet	apple	the	planet	the	apple	the	apple
Doc1	1	0	0	1	1	0	1	1	1	1
Doc2	1	1	0	0	1	1	1	1	1	1
Doc3	1	1	1	1	1	1	1	1	1	1
Doc4	1	0	1	1	1	1	1	1	1	1

Inverted Lists

Conceptually an inverted list looks like an object

apple

df:	4356
docid:	42
tf:	3
locs:	14
	83
	157
docid:	94
:	
:	

Usually it is stored on disk as a sequence of integers

apple

4356
42
3
14
83
157
94
:
:

16

© 2021, Jamie Callan

16

Inverted List Indexes: Compression

Usually inverted lists are compressed – why?

- Save disk space? Favor aggressive compression algorithms
- Save time? Favor simple compression algorithms
 - I/O savings > CPU time required to uncompress

Today, the most common goal is to reduce query time

Algorithms:

- Gap encoding
- Restricted variable-length (RVL) encoding
- The book also covers slower, more aggressive algorithms

17

© 2021, Jamie Callan

17

Inverted File Compression: Delta Gap (“DGap”) Encoding

Store the differences between numbers (“D Gaps”)

- Increases probability of smaller numbers
- A more skewed distribution
- Lower entropy

Stemming also increases the probability of smaller numbers

- Why?

Before

Doc ID	121
TF	3
Loc	18
Loc	47
Loc	68
DocID	135
TF	2
Loc	22
Loc	35

After

Doc ID	121
TF	3
Loc	18
Loc	29
Loc	21
DocID	14
TF	2
Loc	22
Loc	13

18

© 2021, Jamie Callan


18

Inverted File Compression: Variable Byte Encoding

Variable byte encoding stores a number in a sequence of bytes



Each byte contains a flag and 7 bits of payload (the number)



The flag indicates whether this is the last byte in the sequence

0: Not the last byte 1: The last byte

Concatenate the payload bits to reconstruct the number

19

© 2021, Jamie Callan

19

Inverted File Compression: Variable Byte Encoding

Decimal:	5					7 bits
Binary:	00000000	00000000	00000000	00000101		
Compressed:				10000101		

Decimal:	57					7 bits
Binary:	00000000	00000000	00000000	00111001		
Compressed:				10111001		

↑
The flag identifies the last byte

20

© 2021, Jamie Callan

20

Inverted File Compression: Variable Byte Encoding

Decimal:	127				7 bits └───┘
Binary:	00000000	00000000	00000000	01111111	
Compressed:				11111111	

Decimal:	128		7 bits └───┘	7 bits └───┘
Binary:	00000000	00000000	00000000	10000000
Compressed:			10000001	00000000
			Last byte	Byte ₀

Decimal:	131		7 bits └───┘	7 bits └───┘
Binary:	00000000	00000000	00000000	10000011
Compressed:			10000001	00000011

21

© 2021, Jamie Callan

21

Inverted File Compression: Variable Byte Encoding

Decimal:	613,521		7 bits └───┘	7 bits └───┘	7 bits └───┘
Binary:	00000000	00001001	01011100	10010001	
Compressed:		10100101	00111001	00010001	
		Last byte	Byte ₁	Byte ₀	

22

© 2021, Jamie Callan

22

Inverted File Compression: Variable Byte Decoding

Byte₂
Byte₁
Byte₀

Compressed:

10100101
└─┘

00111001
└─┘

00010001
└─┘

Last?
↑
↑
Payload

Initial:

00000000

00000000

00000000

00000000

After Byte₀:

00000000

00000000

00000000

00010001
└─┘

After Byte₁:

00000000

00000000

00011100
└─┘

10010001

After Byte₂:

00000000

00001001
└─┘

01011100

10010001

Decimal:

613,521

23

Inverted File Compression: Variable Byte Encoding

$[0 \dots 2^7-1]:$	1 byte:	1xxxxxxx
$[2^7 \dots 2^{14}-1]:$	2 bytes:	1xxxxxxx0xxxxxxx
$[2^{14} \dots 2^{21}-1]:$	3 bytes:	1xxxxxxx0xxxxxxx0xxxxxxx
:	:	:
:	:	:
:	:	:

Can store numbers of arbitrary size when needed

Advantages:

- Encoding and decoding can be done very efficiently
- Can find the n^{th} number without decoding the previous numbers

24

© 2021, Jamie Callan

24

Inverted File Compression

There are other inverted list compression algorithms

- E.g., Gamma and Delta codes
 - See the textbook for details
- **Note:** DGap Encoding \neq Delta Code

The most effective compression algorithms are ...

- About 15-20% smaller than variable byte encoding
- Slower than restricted variable length encoding

Disks are cheap, and speed is important

- So restricted variable length compression is a common solution

25

© 2021, Jamie Callan

25

Inverted File Compression: Summary

A compressed inverted file, without positional information:

- Less than 10% the size of the original text

A compressed inverted file with positional information:

- 15-20% the size of the original text

26

© 2021, Jamie Callan

26

Lecture Outline

- Building inverted lists on a single processor
- Inverted lists and inverted files
 - Inverted list compression
 - Inverted list optimizations
- Forward indexes
- Index updates

27

© 2021, Jamie Callan

27

Inverted List Optimizations: Multiple Inverted Lists Per Term

Storing several types of inverted list per term improves efficiency

- Binary: For unranked Boolean
- Frequency: For ranking with scores
 - ~2× longer than binary lists
- Positional: For #NEAR, #WINDOW
 - ~2× longer than frequency lists

Use as little data as possible for each task

- Reduced I/O and reduced computation

Cost: Extra disk space

Binary

df:	4356
docid:	42
docid:	94
:	
:	

Frequency

df:	4356
docid:	42
tf:	2
docid:	94
:	
:	

Positional

df:	4356
docid:	42
tf:	2
locs:	14
	83
docid:	94
:	
:	

28

© 2021, Jamie Callan

28

Inverted List Optimizations: Skip Lists

Skip lists are pointers that allow parts of the inverted list to be skipped

One heuristic

- List length: df_t
- A skip pointer every $\sqrt{df_t}$ documents

Purpose

- Reduced computation
- Reduced I/O
 - If inverted lists are read in blocks

Inverted
List
With
Skip
Pointers

```

df  25
ctf 37
skip past doc 19
doc  3, tf 3, locs ...
doc  7, tf 1, locs ...
doc 10, tf 2, locs ...
doc 13, tf 1, locs ...
doc 19, tf 4, locs ...
skip past doc 44
doc 23, tf 1, locs ...
doc 27, tf 2, locs ...
doc 32, tf 1, locs ...
doc 41, tf 1, locs ...
doc 44, tf 1, locs ...
skip past doc 84
doc 57, tf 5, locs ...
:   :   :   :   :
    
```

29

© 2021, Jamie Callan

29

Inverted List Optimizations: Skip Lists

When is a skip list useful? Consider #NEAR/3 (jamie apple)

42 \neq 43, so advance the pointer
with the smaller docid

jamie		apple	
df:	23	df:	1,033,436
→ docid:	42	→ docid:	43
tf:	3	tf:	3
docid:	59,356	docid:	49
tf:	1	tf:	1
:		:	

Document locations are
not shown due to lack
of space on the slide

30

© 2021, Jamie Callan

30

Inverted List Optimizations: Skip Lists

When is a skip list useful? Consider #NEAR/3 (jamie apple)

59,356 \neq 43, so advance the pointer with the smaller docid

- But advancing to the next 'apple' document is inefficient
- Better to advance the 'apple' pointer to at least docid 59,356

jamie		apple	
df:	23	df:	1,033,436
docid:	42	df:	43
tf:	3	tf:	3
docid:	59,356	docid:	49
tf:	1	tf:	1
:		:	

Note: QryIop.java has docIteratorAdvanceTo (docid)

- Advance to docid, or beyond if docid isn't in the list
- It would be easy to add skip lists to the QryEval code

Document locations are not shown due to lack of space on the slide

31

© 2021, Jamie Callan

31

Inverted List Optimizations: Skip Lists

Skip lists are useful for any query operator that needs all of its arguments to occur in a document

- #NEAR, #WINDOW, #SYN
- Boolean AND
- More advanced query operators that we haven't covered

Skipping can also occur when score calculations are complex

- Some query evaluation optimizations stop calculating a document score when it becomes obvious that the score is low

32

© 2021, Jamie Callan

32

Inverted List Optimizations: Top-Docs (Champion) Lists

Main idea

- Some inverted lists are long
- Most queries only need to return ≤ 100 documents
- Why rank all documents if only 100 are needed?

Top-docs lists:

- Truncated inverted lists that contain only the best docs
- Reduced I/O and reduced computation
- Lower recall

“apple” Inverted List

Doc 1
tf 2
Doc 2
tf 4
: : :
Doc 258392
tf 3
Doc 258393
tf 5
: : :
Doc 1025429
tf 6
Doc 1025430
tf 4

“apple” Top-Docs List

Doc 1025429
tf 6
Doc 258393
tf 5
: : :
Doc 2
tf 4
Doc 1025430
tf 4

1,000 documents
1,025,430 documents

33

© 2021, Jamie Callan

33


Inverted List Optimizations: Top-Docs (Champion) Lists

How are top-docs lists constructed?

- Select documents to go in the list by...
 - tf
 - PageRank
 - ...
- Order the top-docs list by document id
 - Faster, if the whole list is read
 - May require multiple lists of different lengths
- Order the top-docs list by tf
 - Supports reading a variable amount of list
 - Requires just one list
- ...

Inverted lists for “apple”

All docs 100 docs 200 docs



34

© 2021, Jamie Callan

34

Inverted List Optimizations: Top-Docs (Champion) Lists

How many terms are frequent enough to have a top-docs list?

- **Linux filesystem page size:** 4096 bytes
 - A page is the minimum unit of I/O
 - Top-docs lists for lists of less than 4096 bytes don't save I/O
- **How many terms have (compressed) inverted lists longer than 4096 bytes?**
 - Terms with $\text{ctf} \geq 1,000$
 - » Probably higher than 1,000, but a careful answer requires corpus analysis

35


© 2021, Jamie Callan

35

Inverted List Optimizations: Top-Docs (Champion) Lists

How many terms are frequent enough to have a top-docs list?

- **Zipf's Law:** $\text{Rank} \times \text{Frequency}_c = A \times N$
- **Rank of a word that occurs once ($\text{ctf}=1$):** $A \times N / 1$
 - Also an estimate of the vocabulary size
- **Rank of a word that occurs 1,000 times:** $A \times N / 1000$
- **The percentage of terms with $\text{ctf} \geq 1000$:**
 $(A \times N / 1000) / (A \times N) = 1 / 1000 = 0.1\%$
- **So ... if the vocabulary is 1,000,000 terms**
 - There are fewer than 1,000 top-docs lists
 - Each list is perhaps 4-8 KB long
 - So ... 4-8 MB



36

© 2021, Jamie Callan

36

Inverted List Optimizations

A high-level view of inverted list optimizations

	Reduces I/O	Reduces Computation	Reduces Accuracy
Compression	++	—	
Multiple inverted lists per term	++	+	
Skip lists	Maybe	++	
Top docs lists	+++	++++	Sometimes

37

© 2021, Jamie Callan

37

Lecture Outline

- Building inverted lists on a single processor
- Inverted lists and inverted files
 - Inverted list compression
 - Inverted list optimizations
- **Forward indexes**
- Index updates

38

© 2021, Jamie Callan

38

Forward Indexes

Suppose I want to know which words occur in documents about Microsoft ... how would I do it?

- This is a common component of text mining tasks
- E.g., sentiment analysis of documents about Microsoft
- E.g., query expansion, relevance feedback

First step: Use an inverted list to find out which documents contain the word Microsoft

- Easy

Now what?

39

© 2021, Jamie Callan

39

Forward Indexes

Sometimes your software needs to know what terms are in the document ... how does it find out?

Parse the document again?

- A little slow (but not terrible, because indexing is fast)
- Done when storage is expensive / small

Store the parsed document?

- Fast
- Done when storage is cheap

40

© 2021, Jamie Callan

40

Forward Indexes

Forward indexes store the indexed form of a document

- The location of every term that made it into the index
 - Term id, location
 - Information about where stopwords appeared
- Optionally: Information about document structure
 - Field names and extents

41

© 2021, Jamie Callan

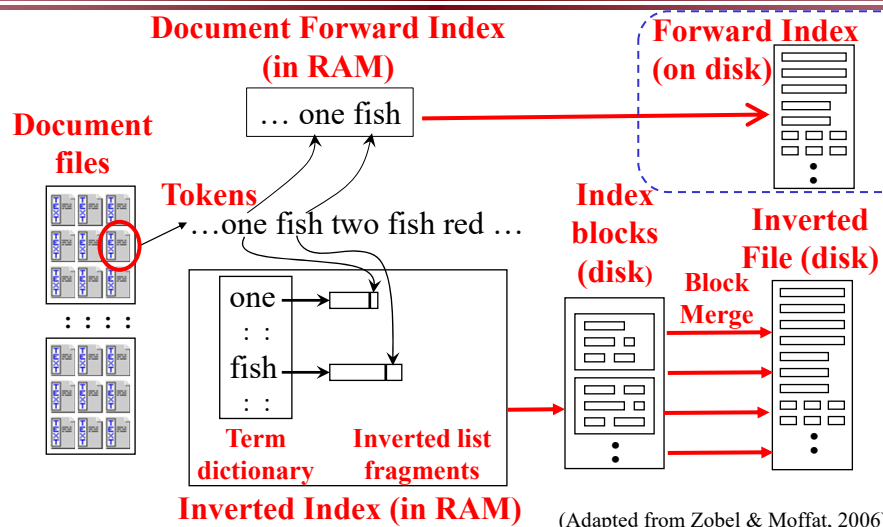
41

How Forward Indexes are Built: Single Processor

Corpus Vocabulary

	one	fish	two	fish	red	...	one
Doc1	1	1	0	1	1	...	1
Doc2	1	1	0	0	1	...	0
Doc3	1	1	1	1	0	...	1
Doc4	1	1	1	1	1	...	1

Forward Index Lists



42

© 2021, Jamie Callan

42

How Does Indri Do It? Two Different Approaches

Indri provides two classes for accessing forward indexes

- **Via the `indri::index` class**
 - A somewhat low-level access to the index
 - Very efficient, not always very friendly
- **Via the `QueryEnvironment` class**
 - Higher-level, more abstract access to the index
 - Somewhat less efficient, somewhat more user friendly

We start with the `indri::index` class because it exposes the data structures more clearly

43

© 2021, Jamie Callan

43

How Does Indri Do It? The `indri::index::TermList` Class

Text: “OBAMA STATE OF THE UNION SPEECH
President Barack Obama delivered ...”

int terms []

obama	41321
state	34127
OOV	0
OOV	0
union	25434
speech	9982
president	98476
barack	12653
obama	41321
deliver	34376
	: :

**Term ids in
the document.
0: stopword**

44

© 2021, Jamie Callan

44

How Does Indri Do It? The indri::index::TermList Class

Text: "OBAMA STATE OF THE UNION SPEECH
President Barack Obama delivered ..."

	int terms []	fields []
obama	41321	int begin;
state	34127	int end;
OOV	0	int id;
OOV	0	
union	25434	0, 5, 3 title
speech	9982	6, 99, 4 body
president	98476	6, 16, 18 sentence
barack	12653	17, 34, 18 sentence
obama	41321	: : : : : :
deliver	34376	
	: :	

**Term ids in
the document.
0: stopword**

45

© 2021, Jamie Callan

45

How Does Indri Do It? The DocumentVector Class

Text: "OBAMA STATE OF THE UNION SPEECH
President Barack Obama delivered ..."

	string stems[]	int positions[]	fields []	int begin; int end; string name;
A list of all terms that occur in the document. On disk, term ids are stored instead of strings.	OOV	1	0, 6, title	
	obama	2	6, 99, body	
	state	0	6, 16, sentence	
	union	0	16, 34, sentence	
	speech	3	: : : :	
	president	4		
	barack	5		
	delivered	6		
	: :	1		
		7		
	:			

**Indexes
into the
stems list**

**Positions of fields.
"Begin" is inclusive
(contained in the field).
"End" is exclusive
(not contained in the field).**

46

© 2021, Jamie Callan

46

Lecture Outline

- Building inverted lists on a single processor
- Inverted lists and inverted files
 - Inverted list compression
 - Inverted list optimizations
- Forward indexes
- Index updates

47

© 2021, Jamie Callan


47

Inverted File Management: Static File (No Updates)

Access
Information
(Small File)



Inverted Lists
(Large File)



- Create files when inverted list fragments are merged
- There is no empty space between inverted lists
- Lists are stored in canonical order (e.g., alphabetic)
- Easy to create, very space efficient
- Very difficult to update; easier to rebuild
 - Update by merging fragments with file to create new file

48

© 2021, Jamie Callan

48

Updating Indexes

Indexes are expensive to update

- Suppose a new document contains 100 unique terms
- Adding that document means updating 100 inverted lists
 - Acquire lock, read list, write list, release lock
 - A lot of complexity, a lot of I/O
- Adding one document is tolerable, adding several is expensive

Updates are often done in batches

- Update every day, or after N documents arrive, or ...
- Parse documents to generate index modifications
- Update each inverted list for all documents in the batch

49

© 2021, Jamie Callan

49

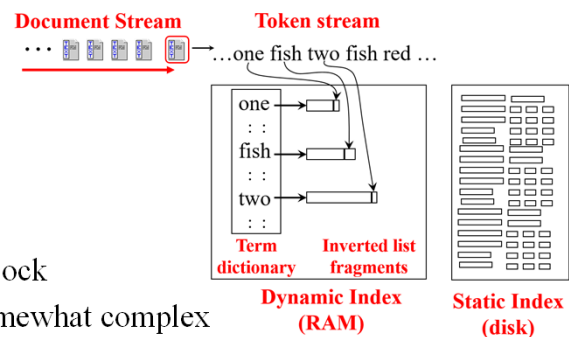
Updating Indexes

Sometimes dynamic updates are unavoidable

- E.g., news, Twitter, ...

Split index into dynamic and static parts

- The dynamic index is small
- The static index is big
- Make updates to the dynamic index
 - Acquire lock, read list, update list, write lock
 - Faster because lists are small, but still somewhat complex
- Search both static (big) and dynamic (small) components
- Periodically merge dynamic into static



50

© 2021, Jamie Callan

50

Deleting Documents

Deleting a document is an expensive operation

- If the document contains N terms, must update N inverted lists
- A major problem in a system that is being used dynamically

Delete lists are a less expensive option

- When a document is deleted, add its id to a delete list
 - Don't actually delete it from the index
- When doing a search
 - Evaluate the query to produce a ranked list
 - Scan the list, removing any documents on the delete list
- When the delete list becomes large
 - Garbage collect the inverted lists, or rebuild the index

51

© 2021, Jamie Callan

51

Lecture Outline

- **Building inverted lists on a single processor**
- **Inverted lists and inverted files**
 - Inverted list compression
 - Inverted list optimizations
- **Forward indexes**
- **Index updates**

52

© 2021, Jamie Callan

52

For More Information

- I.H. Witten, A. Moffat, and T.C. Bell. “Managing Gigabytes.” Morgan Kaufmann. 1999.
- G. Salton. “Automatic Text Processing.” Addison-Wesley. 1989.
- J. Zobel and A. Moffat. “Inverted files for text search engines.” *ACM Computing Surveys*, 38 (2). 2006.